

High-level system constraints:

1. Will the service run on-premise (your hardware) or in a cloud environment (AWS/GCP)?
2. Expected inference load (images per second, concurrency)?
3. Any security / compliance requirements (e.g., no data leaving your network)?

Clarification of existing pipeline (no code, just description):

1. How does your current DINOv3 model receive images and return predictions?
2. What is the exact input/output format (normalized boxes, feature vectors, etc.)?
3. Which heads (BC, ETB, YM, XSA) are used and what are their output classes?

Annotation tool preference:

1. Do you want to continue with Labelbox, switch to Label Studio (open-source), or use a lightweight custom UI?
2. Who will perform the annotations (internal team, outsourced)?

Decision on advanced features:

1. Do you want ensemble uncertainty (3-5 models), hybrid sampling (uncertainty+diversity), pseudo-labeling, or all of them?
2. If ensemble, what is the acceptable inference latency trade-off?